

Assignment 02

Task 1

Loading the UniversalBank.csv and display the summary of dataset

```
data <- read.csv("UniversalBank.csv")
summary(data)
```

```
##           ID           Age           Experience           Income           ZIP.Code
## Min.      : 1      Min.   :23.00      Min.   : -3.0      Min.    : 8.00      Min.    : 9307
## 1st Qu.:1251      1st Qu.:35.00      1st Qu.:10.0      1st Qu.: 39.00      1st Qu.:91911
## Median :2500      Median :45.00      Median :20.0      Median : 64.00      Median :93437
## Mean     :2500      Mean    :45.34      Mean    :20.1      Mean    : 73.77      Mean    :93153
## 3rd Qu.:3750      3rd Qu.:55.00      3rd Qu.:30.0      3rd Qu.: 98.00      3rd Qu.:94608
## Max.     :5000      Max.     :67.00      Max.     :43.0      Max.     :224.00      Max.     :96651
##           Family           CCAvg           Education           Mortgage
## Min.      :1.000      Min.    : 0.000      Min.     :1.000      Min.     : 0.0
## 1st Qu.:1.000      1st Qu.: 0.700      1st Qu.:1.000      1st Qu.: 0.0
## Median :2.000      Median : 1.500      Median :2.000      Median : 0.0
## Mean     :2.396      Mean     : 1.938      Mean     :1.881      Mean     : 56.5
## 3rd Qu.:3.000      3rd Qu.: 2.500      3rd Qu.:3.000      3rd Qu.:101.0
## Max.     :4.000      Max.     :10.000      Max.     :3.000      Max.     :635.0
## Personal.Loan Securities.Account CD.Account           Online
## Min.      :0.000      Min.    :0.0000      Min.     :0.0000      Min.     :0.0000
## 1st Qu.:0.000      1st Qu.:0.0000      1st Qu.:0.0000      1st Qu.:0.0000
## Median :0.000      Median :0.0000      Median :0.0000      Median :1.0000
## Mean     :0.096      Mean     :0.1044      Mean     :0.0604      Mean     :0.5968
## 3rd Qu.:0.000      3rd Qu.:0.0000      3rd Qu.:0.0000      3rd Qu.:1.0000
## Max.     :1.000      Max.     :1.0000      Max.     :1.0000      Max.     :1.0000
##           CreditCard
## Min.      :0.000
## 1st Qu.:0.000
## Median :0.000
## Mean     :0.294
## 3rd Qu.:1.000
## Max.     :1.000
```

Now observe the attribute of dataframe

```
str(data)
```

```
## 'data.frame':    5000 obs. of  14 variables:
## $ ID              : int  1 2 3 4 5 6 7 8 9 10 ...
## $ Age              : int  25 45 39 35 35 37 53 50 35 34 ...
## $ Experience        : int  1 19 15 9 8 13 27 24 10 9 ...
```

```
## $ Income      : int  49 34 11 100 45 29 72 22 81 180 ...
## $ ZIP.Code    : int  91107 90089 94720 94112 91330 92121 91711 93943 90089 93023 ...
## $ Family      : int   4 3 1 1 4 4 2 1 3 1 ...
## $ CCAvg       : num  1.6 1.5 1 2.7 1 0.4 1.5 0.3 0.6 8.9 ...
## $ Education   : int   1 1 1 2 2 2 2 3 2 3 ...
## $ Mortgage    : int   0 0 0 0 0 155 0 0 104 0 ...
## $ Personal.Loan : int   0 0 0 0 0 0 0 0 0 1 ...
## $ Securities.Account: int  1 1 0 0 0 0 0 0 0 0 ...
## $ CD.Account  : int   0 0 0 0 0 0 0 0 0 0 ...
## $ Online      : int   0 0 0 0 0 1 1 0 1 0 ...
## $ CreditCard  : int   0 0 0 0 1 0 0 1 0 0 ...
```

Task 1

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(caret)
```

```
# Load the dataset
```

```
df <- read.csv("UniversalBank.csv", header = TRUE)
```

```
# Create a data frame for the new customer
```

```
new_customer <- data.frame(Age = 40, Experience = 10, Income = 84, Family = 2,
                           CCAvg = 2, Education_1 = 0, Mortgage = 0, Securities.Account = 0,
                           CD.Account = 0, Online = 1, CreditCard = 1)
```

```
# Transform categorical predictors into dummy variables
```

```
df$Education <- as.factor(df$Education)
df$Personal.Loan <- as.factor(df$Personal.Loan)
df$Securities.Account <- as.factor(df$Securities.Account)
df$CD.Account <- as.factor(df$CD.Account)
df$Online <- as.factor(df$Online)
df$CreditCard <- as.factor(df$CreditCard)
```

```
# Partition the dataset into 60% training and 40% validation sets
```

```
set.seed(123)
trainIndex <- createDataPartition(df$Personal.Loan, p = 0.6, list = FALSE)
train_set <- df[trainIndex, ]
valid_set <- df[-trainIndex, ]
```

```
#remove column ID and ZipCode
```

```
library(class)
```

```
# Apply k-NN classification with k = 1
```

```
knn_pred <- knn(train_set[, -c(1, 5,10)], valid_set[, -c(1, 5,10)], train_set$Personal.Loan, k = 1)
```

```
# Predict the new class of new customer
```

```
new_cust_pred <- knn(train_set[, -c(1, 5,10)], new_customer, train_set$Personal.Loan, k = 1)
new_cust_pred
```


```
## [1] 0
## Levels: 0 1
```

task 2 What is the choice of the k that balances between the overfitting and ignoring the predictor information

```
# Apply k fold cross-validation to compute the best value of k
set.seed(123)
n_fold <- trainControl(method = "cv", number = 10)
n_seq <- seq(1, 20, by = 1)
model <- train(Personal.Loan ~ ., data = train_set[, -c(1, 5)], method = "knn",
               trControl = n_fold, tuneGrid = data.frame(k = n_seq))
model
```

```
## k-Nearest Neighbors
##
## 3000 samples
## 11 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 2700, 2700, 2700, 2701, 2700, 2700, ...
## Resampling results across tuning parameters:
##
## k Accuracy Kappa
## 1 0.9056742 0.4196198
## 2 0.8986609 0.3944403
## 3 0.9126609 0.4182487
## 4 0.9093276 0.4045227
## 5 0.9059920 0.3405011
## 6 0.9066598 0.3471799
## 7 0.9083332 0.3358333
## 8 0.9023409 0.2889865
## 9 0.9056698 0.2990069
## 10 0.9053354 0.3018046
## 11 0.9033342 0.2722690
## 12 0.9039976 0.2683909
## 13 0.9039987 0.2519058
## 14 0.9063343 0.2793607
## 15 0.9079998 0.2869012
## 16 0.9049987 0.2757670
## 17 0.9060009 0.2770585
## 18 0.9099976 0.2952039
## 19 0.9063343 0.2579730
## 20 0.9073309 0.2789819
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 3.
```

```
# Plot different values of k
plot(model)
```



```
### Task 3 Show the confusion matrix for validation data that result form using the best k
```

```
pred <- knn(train = train_set[,-10],test = valid_set[,-10], cl = train_set[,10], k=3)
confusionMatrix(pred, valid_set[,10])
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    0    1
```

```
##           0 1780  180
```

```
##           1   28   12
```

```
##
```

```
##           Accuracy : 0.896
```

```
##           95% CI : (0.8818, 0.909)
```

```
##           No Information Rate : 0.904
```

```
##           P-Value [Acc > NIR] : 0.8938
```

```
##
```

```
##           Kappa : 0.0728
```

```
##
```

```
##           McNemar's Test P-Value : <2e-16
```

```
##
```

```
##           Sensitivity : 0.9845
```

```
##           Specificity : 0.0625
```

```
##          Pos Pred Value : 0.9082
##          Neg Pred Value : 0.3000
##          Prevalence     : 0.9040
##          Detection Rate : 0.8900
##          Detection Prevalence : 0.9800
##          Balanced Accuracy : 0.5235
##
##          'Positive' Class : 0
##
```

Task 4

Consider the following customer: Age = 40, Experience = 10, Income = 84, Family = 2, CCAvg = 2, Education_1 = 0, Education_2 = 1, Education_3 = 0, Mortgage = 0, Securities Account = 0, CD Account = 0, Online = 1 and Credit Card = 1. Classify the customer using the best k.

```
new_customer2 <- data.frame(Age = 40, Experience = 10, Income = 84, Family = 2, CCAvg = 2, Education = 0,
knn_build <- knn(train = train_set[, -c(1,5,10)], test = new_customer2, cl = train_set[, 10], k=3)
knn_build
```

```
## [1] 0
## Levels: 0 1
```

Task 5

Repartition the data, this time into training, validation, and test sets (50% : 30% : 20%). Apply the k-NN method with the k chosen above. Compare the confusion matrix of the test set with that of the training and validation sets. Comment on the differences and their reason.

```
## Repartition the data into training, validation, and test sets (50% : 30% : 20%)
set.seed(123) # Setting seed to reproduce results
index_train <- createDataPartition(df$Personal.Loan, p = 0.5, list = FALSE)
train_set <- df[index_train, ]
index_val_test <- createDataPartition(df[-index_train, ]$Personal.Loan, p = 0.6, list = FALSE)
val_test_set <- df[-index_train, ]
val_set <- val_test_set[index_val_test, ]
test_set <- val_test_set[-index_val_test, ]

# Set up the train control with 10-fold cross-validation
trControl <- trainControl(method="cv", number=10)

set.seed(123)
valid_pred <- knn(train_set[, -c(1,5,10)], val_set[, -c(1,5,10)], cl = train_set[, 10], k = 3)
test_pred <- knn(train_set[, -c(1,5,10)], test_set[, -c(1,5,10)], cl = train_set[, 10], k = 3)
train_pred <- knn(train_set[, -c(1,5,10)], train_set[, -c(1,5,10)], cl = train_set[, 10], k = 3)

# Create confusion matrices
confusion_train <- confusionMatrix(train_pred, train_set$Personal.Loan)
confusion_val <- confusionMatrix(valid_pred, val_set$Personal.Loan)
confusion_test <- confusionMatrix(test_pred, test_set$Personal.Loan)
```

```
# Display the confusion matrices
confusion_train
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 2223   91
##           1   37  149
##
##           Accuracy : 0.9488
##           95% CI : (0.9394, 0.9571)
##           No Information Rate : 0.904
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.672
##
## Mcnemar's Test P-Value : 2.805e-06
##
##           Sensitivity : 0.9836
##           Specificity : 0.6208
##           Pos Pred Value : 0.9607
##           Neg Pred Value : 0.8011
##           Prevalence : 0.9040
##           Detection Rate : 0.8892
##           Detection Prevalence : 0.9256
##           Balanced Accuracy : 0.8022
##
##           'Positive' Class : 0
##
```

```
confusion_val
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 1311   87
##           1   45   57
##
##           Accuracy : 0.912
##           95% CI : (0.8965, 0.9259)
##           No Information Rate : 0.904
##           P-Value [Acc > NIR] : 0.1566455
##
##           Kappa : 0.417
##
## Mcnemar's Test P-Value : 0.0003589
##
##           Sensitivity : 0.9668
##           Specificity : 0.3958
##           Pos Pred Value : 0.9378
```

```

##          Neg Pred Value : 0.5588
##          Prevalence : 0.9040
##          Detection Rate : 0.8740
##    Detection Prevalence : 0.9320
##          Balanced Accuracy : 0.6813
##
##          'Positive' Class : 0
##

```

confusion_test

```

## Confusion Matrix and Statistics
##
##          Reference
## Prediction  0   1
##          0 876  59
##          1  28  37
##
##          Accuracy : 0.913
##          95% CI : (0.8938, 0.9297)
##    No Information Rate : 0.904
##    P-Value [Acc > NIR] : 0.181360
##
##          Kappa : 0.4142
##
##    McNemar's Test P-Value : 0.001298
##
##          Sensitivity : 0.9690
##          Specificity : 0.3854
##          Pos Pred Value : 0.9369
##          Neg Pred Value : 0.5692
##          Prevalence : 0.9040
##          Detection Rate : 0.8760
##    Detection Prevalence : 0.9350
##          Balanced Accuracy : 0.6772
##
##          'Positive' Class : 0
##

```

Observation The confusion matrices reveal good model accuracy but highlight issues with specificity, indicating potential difficulties in correctly identifying negative cases. This, combined with a drop in performance on validation and test data compared to training data, suggests possible overfitting and a need for model adjustments or additional strategies to address class imbalance.